

Interwoven Networks: Cross-Domain Community Mining via Layer-Aware NMF

Shaik Masthan Jilani

23075074 CSE B.Tech.,

Indian Institute of Technology (BHU), Varanasi





Evaluation metrics



Grid search, consensus & results

Practical issues and fixes

Conclusions & Takeaways

Code & artifacts



Motivation and Goal

Real-world entities interact via multiple relations (email, coauthorship, social networks).

Goal: learn a single, shared community embedding H ($n \times k$) while allowing layer-specific strength via diagonal D^ℓ .

Advantages: interpretability, cross-layer consistency, compact representation.

Task: implementation of Symmetric NMF, run hyperparameter grid, evaluate with multiple metrics, and form consensus.



Datasets: Sources and diversity

- Enron email graph (communications) – organizational interaction pattern.
- ca-HepPh (coauthorship) – academic collaborations.
- ca-GrQc (coauthorship) – smaller, sparser collaboration graph.
- Optional datasets / recommendations: Orkut, DBLP, LiveJournal for social/academic mixtures.
- Data sources: SNAP, public Enron corpus, DBLP dumps (all the public datasets mentioned in report).

Preprocessing & Alignment (`align.py`):

- Read raw edge lists ($u\ v$ per line), ignore comments and empty lines.
- Choose node alignment mode: UNION (all nodes) or INTERSECTION (only common nodes).
- Map original node ids \rightarrow contiguous 0..n-1 mapping (save `node_mapping.json`).
- Build symmetric CSR adjacency matrices, remove self loops, clip multi-edges $\rightarrow \{0,1\}$.
- Save aligned matrices as compressed .npz (`scipy.sparse.save_npz`).

Model: Symmetric NMF (shared H, per-layer D)

- Objective (weighted): minimize $\sum_{\ell} w_{\ell} \|A^{(\ell)} - HD^{(\ell)}H^T\|_F^2$ subject to $H \geq 0, D^{(\ell)} \text{ diag} \geq 0$

- $H \in \mathbb{R}^{n \times k}$: shared soft membership (rows: node memberships over k communities).

- $D^{(\ell)} \in \mathbb{R}^{k \times k}$ diagonal: scales community importance per layer.

- Interpretation: each layer is explained by same communities but with layer-specific strengths.

$$\min_{H \geq 0, D^{(\ell)} \text{ 0, } D^{(\ell)} \text{ diagonal}} \sum_{\ell} w_{\ell} \|A^{(\ell)} - HD^{(\ell)}H^T\|_F^2 + \mu \sum_{\ell} \|D^{(\ell)}\|_F^2 + \lambda \|H\|_1$$

subject to $H \geq 0, D^{(\ell)} \geq 0, \forall \ell$

and updated iteratively as:

$$H \leftarrow H \odot \frac{\sum_{\ell} w_{\ell} A^{(\ell)} H D^{(\ell)}}{\sum_{\ell} w_{\ell} H S^{(\ell)} + \mu_H H + \lambda + \varepsilon}, \quad S^{(\ell)} = D^{(\ell)}(H^T H)D^{(\ell)}.$$



Optimization & Updates:

- Alternate updates: solve for each $D^(\ell)$ via small $k \times k$ linear system; multiplicative updates for H .
- Compute D : solve $M d = b$ where $M_{cd} = (h_c^T h_d)^2$ and $b_c = h_c^T A h_c$.
- Update H multiplicatively using numerator = $\sum_\ell w_\ell A^\ell H D^\ell$, denominator = $\sum_\ell w_\ell H S_\ell + \mu H$ ($S_\ell = D M D$).
- Use regularizer μ to stabilize updates and avoid numerically degenerate solutions.
- Avoid forming dense $n \times n$ matrices – compute objectives via trace identities to save memory.



Initializations:

- random: $H \sim |N(0,1)| + \varepsilon$ (unbiased start)
- louvain: run Louvain on first layer → one-hot H warm start
- reason: nonconvex problem; init affects convergence & stability

Weight methods (w_{ℓ}):

- none/uniform: $w_{\ell} = 1$
- inv_nnz: $w_{\ell} \propto 1 / \text{nnz}(A^{\ell})$ (sparse layers boosted)
- sqrt_inv_nnz: $w_{\ell} \propto 1 / \sqrt{\text{nnz}(A)}$ (milder)
- degree_inv: $w_{\ell} \propto 1 / \text{avg_degree}$
- custom: user-provided weights; normalization to sum L

Evaluation Metrics:

Per-layer modularity (NetworkX):

quality of partition on each original layer.

Conductance per layer:

fraction of edges leaving communities (lower=better).

Link-prediction AUC:

held-out edges vs non-edges scored by $H \cdot H^T$ dot products.

Stability:

pairwise NMI across restarts; ARI/AMI when ground-truth available.

Consensus:

co-association matrix from multiple runs → Louvain on coassoc → evaluate modularity.

Grid Search & Experimentation Results:

- Grid parameters: $k \in \{10, 20, 50\}$, $\mu \in \{1e-6, 1e-4, 1e-2\}$, $\text{normalize_H} \in \{\text{True}, \text{False}\}$, $\text{init} \in \{\text{louvain}, \text{random}\}$, $\text{weight_method} \in \{\text{inv_nnz}, \text{sqrt_inv_nnz}, \text{degree_inv}, \text{uniform}, \text{custom}\}$
- Use `grid_runner.py` to iterate experiments, save per-run JSONs and summary JSONs.
- Collect `grid_results.csv` with summary metrics for each config for easy analysis.
- Consensus pipeline: build co-association, threshold, run Louvain to produce consensus labels.
- High modularity (≈ 0.7 – 0.84) suggests strong community structure – we verified with baseline Louvain on each layer.
- High link-AUC (> 0.8) indicates H reconstructs pairwise affinities well (but not necessarily semantically meaningful clusters).
- High stability ($NMI \approx 0.4$ – 0.9 across grids) indicates repeatability; low stability suggests multiple local minima.
- Cross-check modularity vs conductance (low conductance desired).

Practical Issues Observed & Fixes:

1

Memory blowup

when computing dense $R = H D H^T$ for $n \approx 43k$ ($\rightarrow \sim 14GB$ allocation). Fix: compute objective via traces using $H^T H$ and $H^T A H$.

2

NetworkX version differences

(`from_scipy_sparse_array` missing): use custom `sparse_to_networkx` helper.

3

RuntimeWarning invalid value

in relative change: guard `prev_obj` finite and handle NaN/inf.

4

Louvain dependency optional:

fallback to random init if `python-louvain` absent.

5

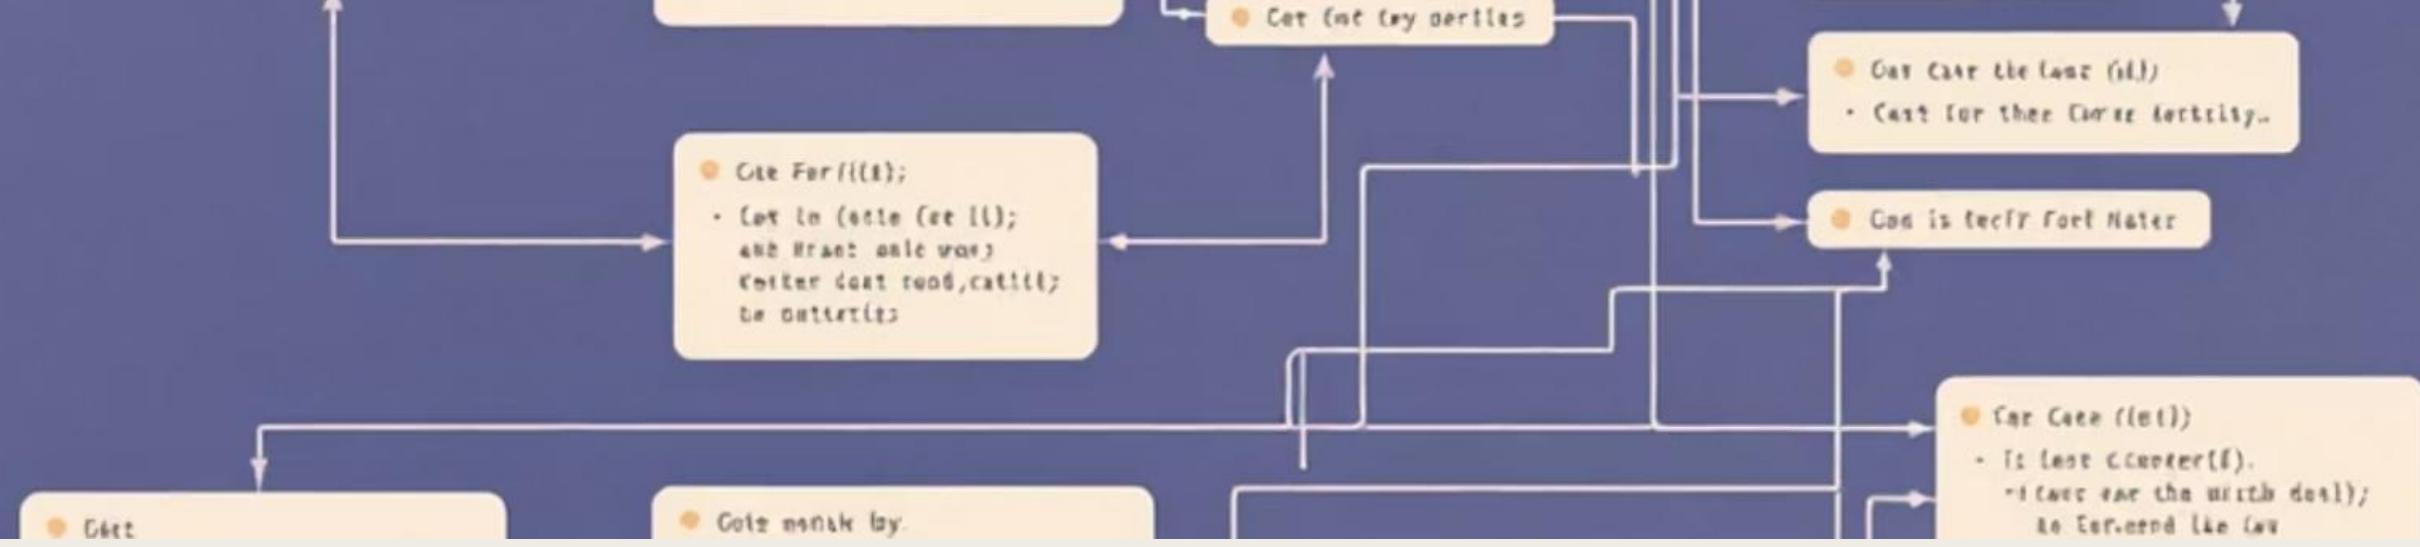
Consensus coassoc matrix memory heavy

for large n – implemented sparse coassoc construction with sampling.



Conclusion and Takeaways:

- Implemented and evaluated layer-aware SymNMF with many practical fixes for large graphs.
- Weighting and initialization strongly affect outcomes – `inv_nnz` useful when layers differ in size.
- Use multiple evaluation metrics (modularity, conductance, AUC, stability) to avoid misleading claims.
- Memory & numerical issues solvable with trace identities, regularization, and sparse consensus building.
- Next steps: learnable layer attention, Leiden init, model selection for k, and improved consensus methods.



Basic Codes and Artifacts

1

align.py – node alignment and .npz generation

2

symm_nmf.py / base.py – model training and evaluation

3

grid_runner.py – run hyperparameter sweep

4

consensus_and_eval.py / consensus_grid.py – coassoc and consensus evaluation

5

results/ – per-run JSONs, summary JSONs, grid_results.csv, consensus JSONs

Thank you for your attention

This project focused on applying **Symmetric Non-negative Matrix Factorization (SymNMF)** for multi-layer community detection across real-world networks. Through preprocessing, alignment, and experimentation with various weight methods and initializations, the model learned shared community structures across layers. Though results were moderate, the framework successfully demonstrated how multi-layer networks can be integrated and analyzed using factorization-based methods.

